

# Practical Work

Pham Huu Minh - 23BI14302

2025-12-05

## Goal of the Practical Work

The objective of this practical work is to implement a one-to-one file transfer system using the Message Passing Interface (MPI). This involves upgrading the existing TCP file transfer implementation to utilize MPI's communication primitives for distributed file transfer between processes.

## Specific Objectives

- Install and configure an MPI implementation (OpenMPI)
- Design an MPI-based architecture for file transfer
- Implement the file transfer mechanism using MPI communication primitives
- Compare the MPI approach with previous TCP Socket and RPC implementations
- Test and verify the correctness of the file transfer system

## Environment Setup

The development environment used for this practical work is **Kali Linux**. Kali Linux provides a stable platform with excellent support for development tools and MPI implementations.

## OpenMPI Installation

**MPI Implementation Used:** OpenMPI

OpenMPI was chosen for its wide availability, excellent documentation, and cross-platform support. The installation was performed using the following commands:

```
sudo apt update  
sudo apt install openmpi-bin openmpi-common libopenmpi-dev
```

To verify the installation, the following commands were used:

```
mpicc -v  
mpirun --version
```

The output confirmed that OpenMPI was successfully installed and ready for use.

## MPI Design for File Transfer

### Rank-based Architecture

The MPI file transfer system uses a two-process model with distinct roles:

- **Rank 0 (Server/Sender):**
  - Reads the input file from disk
  - Determines the file size
  - Sends the file size and file content to Rank 1
- **Rank 1 (Client/Receiver):**
  - Receives the file size from Rank 0
  - Allocates memory for the file data
  - Receives the file content from Rank 0
  - Writes the received data to the output file

#### **Command-line Arguments:**

- `input_file`: Path to the source file to be transferred (used by Rank 0)
- `output_file`: Path where the received file will be saved (used by Rank 1)

## **Message Format**

The communication protocol uses two distinct message tags to differentiate between file size and file data:

- `MSG_TAG_SIZE` (value: 100): Used for sending/receiving the file size as an `MPI_INT`
- `MSG_TAG_FILE` (value: 200): Used for sending/receiving the file content as `MPI_BYTE`

The protocol follows a two-step process:

1. First, Rank 0 sends the file size (integer) with tag `MSG_TAG_SIZE`
2. Then, Rank 0 sends the file data (byte array) with tag `MSG_TAG_FILE`
3. Rank 1 receives both messages in the same order

## **System Organization**

The system consists of a single source file: `mpi_file_transfer.c`

The behavior of the program is determined by the MPI rank:

- When `MPI_Comm_rank()` returns 0, the process executes the sender logic
- When `MPI_Comm_rank()` returns 1, the process executes the receiver logic

The code is organized into modular functions:

- `read_file_to_buffer()`: Handles file reading operations
- `write_buffer_to_file()`: Handles file writing operations
- `sender_process()`: Implements the sender logic (Rank 0)
- `receiver_process()`: Implements the receiver logic (Rank 1)
- `main()`: Initializes MPI, validates arguments, and routes to appropriate process function

# Implementation

## Main Structure and Argument Checking

The main function includes necessary headers, defines constants, and performs initialization:

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <errno.h>

#define MAX_FILE_SIZE 1048576
#define MSG_TAG_SIZE 100
#define MSG_TAG_FILE 200

int main(int argc, char *argv[])
{
    int process_rank, total_processes;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &process_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &total_processes);

    if (total_processes != 2) {
        if (process_rank == 0) {
            fprintf(stderr, "Error: This program requires exactly 2 MPI pr
            fprintf(stderr, "Usage: mpirun -np 2 %s <input_file> <output_f
        }
        MPI_Finalize();
        return EXIT_FAILURE;
    }

    if (argc != 3) {
        if (process_rank == 0) {
            fprintf(stderr, "Usage: %s <input_file> <output_file>\n", argv
        }
        MPI_Finalize();
        return EXIT_FAILURE;
    }

    if (process_rank == 0) {
        sender_process(argv[1]);
    } else {
        receiver_process(argv[2]);
    }

    MPI_Finalize();
    return EXIT_SUCCESS;
}
```

## Server Side: Rank 0

The sender process (Rank 0) reads the file and transmits it to the receiver:

```

static void sender_process(const char *input_filename)
{
    unsigned char *file_data;
    int file_size;

    if (read_file_to_buffer(input_filename, &file_data, &file_size) != 0)
        MPI_Abort(MPI_COMM_WORLD, EXIT_FAILURE);
    return;
}

printf("[Sender] Read %d bytes from '%s', transmitting to receiver...\n",
       file_size, input_filename);

MPI_Send(&file_size, 1, MPI_INT, 1, MSG_TAG_SIZE, MPI_COMM_WORLD);
MPI_Send(file_data, file_size, MPI_BYTE, 1, MSG_TAG_FILE, MPI_COMM_WORLD);

printf("[Sender] Transmission completed successfully.\n");
free(file_data);
}

```



The sender process:

1. Opens the input file and reads its entire content into memory
2. Validates the file size (must be within the 1 MB limit)
3. Sends the file size using `MPI_Send()` with tag `MSG_TAG_SIZE`
4. Sends the file data using `MPI_Send()` with tag `MSG_TAG_FILE`
5. Frees allocated memory and completes

## Client Side: Rank 1

The receiver process (Rank 1) receives the file and writes it to disk:

```

static void receiver_process(const char *output_filename)
{
    int file_size;
    unsigned char *file_data;
    MPI_Status status;

    MPI_Recv(&file_size, 1, MPI_INT, 0, MSG_TAG_SIZE, MPI_COMM_WORLD, &status);

    if (file_size <= 0 || file_size > MAX_FILE_SIZE) {
        fprintf(stderr, "[Receiver] Invalid file size received: %d\n", file_size);
        MPI_Abort(MPI_COMM_WORLD, EXIT_FAILURE);
        return;
    }

    file_data = (unsigned char *)malloc(file_size);
    if (file_data == NULL) {
        perror("[Receiver] Memory allocation failed");
        MPI_Abort(MPI_COMM_WORLD, EXIT_FAILURE);
        return;
    }

    MPI_Recv(file_data, file_size, MPI_BYTE, 0, MSG_TAG_FILE, MPI_COMM_WORLD, &status);

    if (write_buffer_to_file(output_filename, file_data, file_size) != 0)
        free(file_data);
    MPI_Abort(MPI_COMM_WORLD, EXIT_FAILURE);
    return;
}

printf("[Receiver] Successfully received %d bytes and saved to '%s'\n"
       "      file_size, output_filename);
free(file_data);
}

```

The receiver process:

1. Receives the file size using `MPI_Recv()` with tag `MSG_TAG_SIZE`
2. Validates the received file size
3. Allocates memory for the file data
4. Receives the file data using `MPI_Recv()` with tag `MSG_TAG_FILE`
5. Writes the received data to the output file using `fwrite()`
6. Frees allocated memory and completes

## Build Commands

The compilation command used to build the MPI file transfer program:

```
mpicc -Wall -g -o mpi_file_transfer mpi_file_transfer.c
```

Where:

- `mpicc`: OpenMPI's C compiler wrapper
- `-Wall`: Enable all compiler warnings
- `-g`: Include debugging information
- `-o mpi_file_transfer`: Specify the output executable name
- `mpi_file_transfer.c`: Source file

# Execution and Test Results

## Creating a Test File

A test file was created with dummy data for testing:

```
echo "This is a test file for MPI file transfer." > test_mpi.txt
```

Additional test files of various sizes were also created to test the system with different file sizes.

## Running the MPI Program

The MPI program was executed using the following command:

```
mpirun -np 2 ./mpi_file_transfer test_mpi.txt test_mpi_copy.txt
```

### Output logs:

```
[Sender] Read 45 bytes from 'test_mpi.txt', transmitting to receiver...
[Sender] Transmission completed successfully.
[Receiver] Successfully received 45 bytes and saved to 'test_mpi_copy.txt'
```

The output shows that:

- Rank 0 successfully read 45 bytes from the input file
- Rank 0 completed sending the data
- Rank 1 successfully received 45 bytes and wrote them to the output file

## Verifying the Result

To verify that the files are identical, the `diff` command was used:

```
diff test_mpi.txt test_mpi_copy.txt
```

No output from `diff` confirms that the files are identical. Alternatively, the files can be compared using:

```
cat test_mpi.txt
cat test_mpi_copy.txt
```

Both commands show identical content, confirming the successful file transfer.

# Discussion

## MPI vs. TCP Socket Implementation

### Advantages of MPI:

- **Simplified Communication:** MPI provides high-level communication primitives (`MPI_Send`/`MPI_Recv`) that abstract away low-level socket programming details
- **Process Management:** MPI handles process creation, synchronization, and cleanup automatically
- **Portability:** MPI code is portable across different platforms and network configurations
- **Performance:** MPI implementations are highly optimized and can utilize high-

performance interconnects (InfiniBand, etc.)

- **No Manual Network Programming:** No need to handle socket creation, binding, listening, or connection management

#### **Disadvantages of MPI:**

- **Requires MPI Runtime:** Both sender and receiver must have MPI installed and running
- **Less Flexibility:** Less control over low-level network behavior compared to raw sockets
- **Learning Curve:** Requires understanding of MPI concepts (ranks, communicators, tags)

## **MPI vs. RPC Implementation**

#### **Advantages of MPI:**

- **Direct Communication:** MPI provides direct point-to-point communication without the overhead of RPC marshalling/unmarshalling
- **Simpler Protocol:** No need to define interface specifications (like XDR files in RPC)
- **Better Performance:** Lower latency for simple data transfers without RPC overhead
- **No Code Generation:** No need for code generation tools like rpcgen

#### **Disadvantages of MPI:**

- **Less Abstraction:** RPC provides a more abstract interface (function calls) compared to MPI's message-passing model
- **No Interface Definition:** RPC's interface definition files provide better documentation and type safety
- **Platform Dependency:** MPI is primarily used in HPC environments, while RPC is more general-purpose

## **Personal Work**

All tasks for this practical work were completed individually by Pham Huu Minh (23BI14302):

- Research and selection of MPI implementation (OpenMPI)
- Installation and configuration of OpenMPI on Kali Linux
- Design of the MPI-based file transfer architecture
- Implementation of the sender process (Rank 0) logic
- Implementation of the receiver process (Rank 1) logic
- Implementation of file I/O helper functions
- Code compilation and debugging
- Testing with various file sizes and types
- Comparison analysis with TCP Socket and RPC implementations
- Documentation and report writing

## Conclusion

This practical work successfully demonstrated the implementation of a file transfer system using the Message Passing Interface (MPI). The system was able to transfer files between two MPI processes using OpenMPI, with Rank 0 acting as the sender and Rank 1 as the receiver.

The implementation showed that MPI provides a clean and efficient way to implement distributed file transfer, with advantages in terms of code simplicity and performance compared to low-level socket programming. However, it requires the MPI runtime environment and has a different programming model compared to RPC.

The lab successfully demonstrated three different approaches to distributed file transfer:

- **TCP Sockets:** Low-level, flexible, but requires manual network programming
- **RPC:** High-level abstraction with interface definitions, but with marshalling overhead
- **MPI:** Efficient message-passing model, ideal for HPC environments, but requires MPI runtime

Each approach has its own strengths and is suitable for different use cases and environments.